Results and Observations

Grounded evidence from the recorded five-run consistency validation

1. Evaluation Scope

The results below are limited to the recorded navigation-meeting fixture and the five fresh provider analyses stored in **data/validation/consistency/results.json**. Each run used the same transcript, schema, analysis and prompt versions, and model configuration. The fixture contains seven timestamped segments and a 47-second meeting.

Reproducibility boundary: these are measured fixture-level results. They do not establish corpus-wide accuracy, universal LLM determinism, or a general hallucination rate.

2. Re-analysis Stability

Category	Run 1	Run 2	Run 3	Run 4	Run 5	Min. Jaccard
Action items	2	2	2	2	2	100%
Decisions	1	1	1	1	1	100%
Questions	1	1	1	1	1	100%
Follow-ups	0	0	0	0	0	100%

All five runs returned exactly 2 action items, 1 decision, 1 question, and 0 follow-ups. The stored validation report records 100% minimum Jaccard overlap for actions, decisions, and questions against Run 1. Follow-ups were absent from the fixture, so no follow-up extraction accuracy can be inferred.

3. Pipeline Retention and Evidence Matching

Category	Generated	Parsed	Validated	Evidence Matched	Deduplicated
Action items	2	2	2	2	2
Decisions	1	1	1	1	1
Questions	1	1	1	1	1
Follow-ups	0	0	0	0	0

In Run 1, every generated item survived parsing, validation, verbatim evidence matching, and deduplication: 2/2 actions, 1/1 decision, and 1/1 question. The same complete retention pattern is recorded for Runs 2-5.

4. Observations

Observation 1 - Stable core identities. The same source-anchored items were retained across the five runs: actions at segments 3 and 6, the decision at segment 2, and the question at segment 5 (zero-based segment numbering).

Observation 2 - Deterministic reconciliation. Run 1 created 2 actions, 1 decision, and 1 question. Runs 2-5 reported these same items as stable, with no removals and no newly introduced core items. Topic labels varied between runs (2 to 5), showing that the current stability result applies to reconciled core items, not every generated descriptive field.

Observation 3 - Grounding is enforced at the pipeline boundary. The recorded evidence-matched counts equal the generated counts for all extracted core categories in every run. This supports a fixture-level evidence-match rate of 100% (20 of 20 category-stage items across five runs), not a claim that arbitrary model outputs cannot hallucinate.

Observation 4 - Empty categories require cautious interpretation. No follow-up was present in the fixture. The repeated zero therefore measures fixture content, not follow-up recall.

5. Configuration and Traceability

Analysis version: **meeting-intelligence-v3**; prompt version: **exhaustive-evidence-v1**; model: **openai/gpt-oss-120b**; temperature: **0**; top-p: **1**; response format: **json_schema**. Input hash:

`1615c5fecfbb465c9c2d0fe3b8c733743864902fb0f559bd67ca32a9417cd225`.

The manuscript's discourse-classification scores, AMI/ICSI corpus claims, ROUGE values, latency comparisons, and the 42.3% zero-shot baseline are not reproduced by this validation artifact. They should remain provisional until their datasets, predictions, references, and measurement scripts are supplied.

Source artifact generated from recorded validation data. Validation timestamp in the artifact: 2026-09-29T19:48:18.181726+00:00.